Comprehensive Reviewer Report (5 Iterations)

On: "Political Polarization and Market Reactions to Auditor-Change Disclosures"

Review prepared for the author

April 3, 2026

Contents

Executive Summary **2**

1 Iteration 1: Novelty, Contributions, and Overall Assessment **3**

1.1 Novelty Assessment . . . 3

1.2 Evaluation of the Three Claimed Contributions . . . 3

1.3 Abstract and Framing . . . 4

2 Iteration 2: Theory–Empirics Alignment **5**

2.1 Primary vs. Secondary Outcome Inconsistency . . . 5

2.2 The Ambiguity Prediction Failure (P3) . . . 5

2.3 Concavity and the Link to $|CAR|$ . . . 6

2.4 The Event-Type Prediction: Auxiliary or Core? . . . 7

3 Iteration 3: Proof Verification **7**

3.1 Proposition 1 (Polarization and Belief Dispersion) . . . 7

3.2 Propositions A1–A3: A Mathematical Error . . . 8

3.3 Concavity Claim in Section 3.2 . . . 9

3.4 Remark on Proposition A3 Nesting . . . 10

4 Iteration 4: Empirical Design and Identification **10**

4.1 Is Competitiveness Actually Polarization? . . . 10

4.2 Weakness of Alternative Proxies . . . 10

4.3 Sample Construction and Attrition . . . 11

4.4 Forward-Filling and Time Variation . . . 12

4.5 Measurement Error and Classical Attenuation . . . 12

4.6 Endogeneity of Headquarters Location . . . 13

4.7 The Placebo Test Interpretation . . . 13

5 Iteration 5: Redundancy, Notation, and Exposition **13**
5.1 Repeated Passages and Expressions 13
5.2 Notation Consistency 14
5.3 Structural and Expositional Suggestions 15
5.4 Minor Writing Issues 15

6 Summary of Actionable Recommendations **16**

Executive Summary

This report presents five progressively deeper review iterations of the manuscript, covering (1) novelty and contribution assessment, (2) theory–empirics alignment, (3) proof verification, (4) empirical design and identification, and (5) redundancy, notation, and exposition. The paper investigates a genuinely interesting question—whether political polarization shapes how investors interpret mandatory auditor-change disclosures—and the empirical setting (standardized Form 8-K Item 4.01 filings) is well-chosen. The main baseline result (polarization $\rightarrow$ larger $|CAR|$ and $AbVol$) is plausible and reasonably well-supported.

However, five categories of concern are identified:

1. **A mathematical error in the Appendix proofs** (Propositions A1–A3): the claim that $|\sigma(X_D) - \sigma(X_R)|$ is strictly increasing in $|X_D - X_R|$ for arbitrary movements is false due to the saturation of the logistic function. This invalidates the stated results as written and requires either corrected statements or additional regularity conditions (Section 3).

2. **The ambiguity prediction (P3) fails for the primary outcome**: the model's most distinctive cross-sectional prediction—that polarization effects are amplified for high-ambiguity events—is contradicted by the volume results (low-ambiguity events drive the result), and the $|CAR|$ pattern is underpowered. This significantly weakens the paper's second claimed contribution (Section 2).

3. **Inconsistent treatment of primary vs. secondary outcomes**: the theory designates $AbVol$ as primary and $|CAR|$ as secondary, but the abstract, tables, and cross-sectional tests foreground $|CAR|$, creating a presentational tension that an alert referee will notice (Section 2).

4. **The competitiveness measure conflates competitive elections with polarized electorates**: a 50–50 state may contain many moderates rather than two entrenched partisan camps, and the paper does not adequately address this distinction (Section 4).

5. **Substantial redundancy and expositional bloat**: several passages (proxy caveats, ambiguity discussion, channel-agnosticism) are repeated 3–5 times across sections, inflating length without adding content (Section 5).

Each issue is accompanied by specific, actionable suggestions for revision.

1 Iteration 1: Novelty, Contributions, and Overall Assessment

1.1 Novelty Assessment

The paper sits at the intersection of two well-established literatures: (a) political polarization and capital markets [surveyed by][]kempf2024, and (b) auditor-change event studies. The novelty lies in the *intersection*—no prior paper has examined whether polarization shapes market reactions to auditor-change disclosures specifically. This is a legitimate contribution to the auditing literature, which has not absorbed the political-economy-of-finance findings.

However, the underlying mechanism—heterogeneous priors → posterior dispersion → trading volume—is entirely standard from kandel1995 and harris1993. The theoretical contribution is the application of this existing apparatus to a new context, not the development of new theoretical insights. The paper should be more forthright about this: the model formalizes a known mechanism in a new setting rather than deriving surprising predictions.

Major Issue

The paper's novelty is primarily *empirical*—showing that a known mechanism (heterogeneous-prior-driven trading) operates in a specific new setting (auditor-change disclosures). The theory section gives the impression of a deeper theoretical contribution than is warranted.

Suggestion: Shorten the theory section by approximately 30%, framing it explicitly as "applying the Kandel–Pearson framework to the auditor-change setting" rather than developing a new model. Alternatively, if the theory is meant to be a standalone contribution, the authors need to identify what is *new* in the theoretical results beyond applying Bayes' rule to a two-group setup.

1.2 Evaluation of the Three Claimed Contributions

Contribution 1 (polarization → audit disclosure interpretation): This is the most secure contribution. The baseline result ($\beta = 0.005$, $p = 0.040$ for $|CAR|$; $\beta = 0.111$, $p = 0.010$ for $AbVol$) is statistically significant and economically meaningful, and the permutation test ($p = 0.030$) provides additional reassurance.

Contribution 2 (effects concentrated in ambiguous events, distinguishing interpretation from salience): This is the most problematic contribution. The ambiguity test (Table 4) produces results that *contradict* the model's prediction P3 for the primary outcome ($AbVol$): the effect is driven by *low*-ambiguity events ($\beta = 0.114$, $p = 0.008$), not high-ambiguity events ($p = 0.422$). The $|CAR|$ pattern is directionally consistent but insignificant ($p = 0.110$, $N = 120$). The paper acknowledges the volume reversal and labels its explanation a "post-hoc rationalization." A referee targeting the second contribution will find this damaging.

Critical Issue

The paper's second contribution—that polarization effects are concentrated in ambiguous events—is not supported by the data for the theoretically primary outcome. The event-type split (Table 3) partially compensates, but the formal ambiguity test (Table 4) fails on its own terms.

Suggestion: Either (a) reframe the second contribution around the event-type result (non-dismissal vs. dismissal), which is cleaner and significant, or (b) develop a more rigorous theoretical account of *why* low-ambiguity events should drive volume disagreement (this would need to be more than a post-hoc paragraph), or (c) downgrade the ambiguity test from "primary mechanism test" to "exploratory," and be transparent that the prediction is only partially supported.

Contribution 3 (disclosure effectiveness depends on the interpretive environment): This is better characterized as a *policy implication* than a *contribution*. The finding that polarization amplifies market reactions is the empirical contribution; the inference that "disclosure-based regulation has limits" is an interpretation of that finding. Presenting it as a separate contribution inflates the paper's claims.

Minor Issue

Suggestion: Recast the third contribution as a policy implication of the first two, rather than a standalone contribution. This would be both more honest and more persuasive.

1.3 Abstract and Framing

The abstract states that "effects are concentrated in events not initiated by the firm." This refers to the $|CAR|$ result from Table 3. The abstract does not mention that the ambiguity test fails for volume, nor that the DW-NOMINATE proxies are weak. This selective emphasis creates an overly favorable impression.

Major Issue

Suggestion: The abstract should qualify the cross-sectional evidence. For instance: "Cross-sectional tests provide partial support for an interpretation channel: effects are larger for non-dismissal events, though the formal ambiguity interaction does not reach significance for all outcomes."

2 Iteration 2: Theory–Empirics Alignment

2.1 Primary vs. Secondary Outcome Inconsistency

The theory section carefully and correctly argues that $AbVol$ is the "primary formal prediction" of Proposition 1 (via Equation 2), and that $|CAR|$ is a "secondary and less formally established implication" (Section 3.3, lines 412–424). This ordering is theoretically sound: disagreement models predict volume more directly than unsigned returns.

However, the paper's *presentation* inverts this ordering almost everywhere:

- The abstract leads with the $|CAR|$ result ($p = 0.040$) before the $AbVol$ result ($p = 0.010$).
- Tables present $|CAR|$ in the first columns and $AbVol$ in later columns.
- The event-type split (Table 3) and ambiguity test (Table 4) are discussed primarily through the lens of $|CAR|$.
- The introduction reports the event-type split coefficients for $|CAR|$ ($\beta = 0.012$, $p = 0.043$ vs. $\beta = 0.003$, $p = 0.203$) without reporting the corresponding $AbVol$ split.

This creates the impression that the paper selectively foregrounds whichever outcome produces the more favorable result in each test: $AbVol$ when the baseline is at issue (stronger p-value), $|CAR|$ when the event-type split is at issue (significant asymmetry).

Critical Issue

The inconsistency between the theoretical ordering (volume primary, returns secondary) and the presentational ordering (returns first everywhere) will be noticed by referees. If the theory designates $AbVol$ as primary, the empirical presentation must respect that ordering.

Suggestion: (a) Present $AbVol$ before $|CAR|$ in all tables and in the abstract. (b) Report both the $AbVol$ and $|CAR|$ event-type splits in the introduction. (c) If the $AbVol$ event-type split is weaker than the $|CAR|$ split, acknowledge this and discuss why.

2.2 The Ambiguity Prediction Failure (P3)

The model's most distinctive empirical prediction is P3: polarization effects should be larger for high-ambiguity events, where "the objective log-likelihood ratio ℓ is small" (Section 3.4). The paper constructs a `high_ambiguity` indicator (120 of 678 events) and splits the sample.

Results:

- $|CAR|$: High-ambiguity $\beta = 0.020$ ($p = 0.110$); low-ambiguity $\beta = 0.003$ ($p = 0.209$). *Directionally consistent* but insignificant.
- $AbVol$: Low-ambiguity $\beta = 0.114$ ($p = 0.008$); high-ambiguity coefficient is imprecise ($p = 0.422$). *Opposite to the prediction.*

The paper's explanation—that low-ambiguity events "give investors with opposing political priors concrete grounds on which to disagree about severity"—directly contradicts the model, which predicts that when ℓ is large (low ambiguity), the common likelihood ratio dominates and posteriors converge. If investors disagree more about *low-ambiguity* events, the model's account of the ambiguity channel is wrong, not merely imprecisely estimated.

Critical Issue

The volume result for the ambiguity split is not merely a power failure; it is a point estimate that runs in the opposite direction from P3. The post-hoc explanation is internally inconsistent with the formal model. This is the paper's most significant theory–empirics gap.

Suggestions:

1. Develop a formal account (even a brief one) of why identifiable quality signals might *increase* rather than decrease disagreement-driven trading. One possibility: when ℓ is large, all investors update significantly, but the *direction* of the update depends on how investors interpret the quality signal ("the firm upgraded to Big 4 because governance is improving" vs. "the firm was forced to upgrade by regulators"), so disagreement shifts from the magnitude of updating to its interpretation. This would be a different model from the one in the paper.

2. Alternatively, report the ambiguity test with full transparency and explicitly state that P3 is not supported for the primary outcome. Do not label the explanation a "post-hoc rationalization" only in passing; engage directly with what the failure implies for the theory.

3. Consider whether the `high_ambiguity` indicator actually captures the theoretical object ℓ. The indicator combines "no disclosed disagreement" and "lateral quality move." These proxy for filing-level informational content, but they may not map cleanly to the theoretical notion of signal informativeness (p_H/p_L).

2.3 Concavity and the Link to $|CAR|$

The paper correctly notes that $\phi(\pi)$ is concave (Section 3.2), so a mean-preserving spread of priors shifts the average posterior downward. It then states that "the link from posterior variance to $|CAR|$ depends additionally on demand asymmetry, short-sale constraints, or limits to arbitrage that are not explicitly modeled here" (Section 3.3). This is an honest caveat, but the paper then proceeds to test $|CAR|$ as though the theoretical prediction were as strong as for $AbVol$.

Major Issue

Suggestion: Add a brief paragraph either in the theory section or in the hypothesis development explicitly stating what additional assumptions would be needed for the $|CAR|$ prediction to follow. Even a simple demand-based model (e.g., two risk-neutral groups with different posterior means generating a net demand imbalance) would tighten the link and make the $|CAR|$ tests interpretable. Without this, the $|CAR|$ regressions are atheoretical.

2.4 The Event-Type Prediction: Auxiliary or Core?

Section 4 develops the event-type prediction (non-dismissal vs. dismissal) over approximately 40 lines, citing beneish2005, boone2001, and ghosh2015, and carefully distinguishing "event-level signal quality" from "per-filing informational content." This is a substantial theoretical development. But the paper then labels the result "an auxiliary empirical exercise (Table **??**typetab:event_typeher than a core hypothesis." The tension is that the event-type split actually *works better* than the formal ambiguity test (Table 4), yet the paper treats it as secondary.

Major Issue

Suggestion: Either (a) promote the event-type prediction to a formal hypothesis and present Table 3 as a core test, or (b) shorten the theoretical development in Section 4 to match its "auxiliary" status. The current structure devotes significant space to a prediction that is then downgraded, which is confusing.

3 Iteration 3: Proof Verification

3.1 Proposition 1 (Polarization and Belief Dispersion)

The proof in Appendix A is correct.

Step 1 (Variance formula): The two-point distribution calculation is standard and the algebra checks out. The key identities $\phi(\pi_D) - \bar{\phi} = \omega_R(\phi(\pi_D) - \phi(\pi_R))$ and $\phi(\pi_R) - \bar{\phi} = -\omega_D(\phi(\pi_D) - \phi(\pi_R))$ are correct.

Step 2 (Monotonicity): The mean-preserving spread construction $\pi'_D = \pi_D + \omega_R\epsilon$, $\pi'_R = \pi_R - \omega_D\epsilon$ correctly preserves the mean. Strict monotonicity of ϕ then yields $\phi(\pi'_D) > \phi(\pi_D)$ and $\phi(\pi'_R) < \phi(\pi_R)$, so the gap $\phi(\pi'_D) - \phi(\pi'_R) > \phi(\pi_D) - \phi(\pi_R) > 0$. The variance formula gives the result. **No error.**

Remark on mean-preservation: The claim that mean-preservation is sufficient but not necessary is correct. The variance depends only on $|\phi(\pi_D) - \phi(\pi_R)|$, which is increasing in $|\pi_D - \pi_R|$ because ϕ is strictly increasing on $(0, 1)$ (a bounded domain). **No error.**

3.2 Propositions A1–A3: A Mathematical Error

Proposition A1 (Institutional trust channel) states: "Any increase in $|\tau_D - \tau_R|$ strictly increases $\mathrm{Var}(m(\tau_g))$." The proof claims: "Strict monotonicity of m implies $|m(\tau_D) - m(\tau_R)|$ is strictly increasing in $|\tau_D - \tau_R|$."

This claim is false as a general mathematical statement. For a strictly increasing function $f : \mathbb{R} \to \mathbb{R}$, it is *not* true that $|f(a) - f(b)|$ is a strictly increasing function of $|a - b|$ without further constraints on how a and b move.

Counterexample. Let $\sigma(x) = 1/(1 + e^{-x})$ (the logistic function used in the paper). Consider:

$$\begin{aligned}
&\tau_D = 1,\ \tau_R = 0,\ |\tau_D - \tau_R| = 1.\\
&\text{With } a(\pi) = 0,\ \ell = 1: \quad m(1) = \sigma(1) \approx 0.731, \quad m(0) = \sigma(0) = 0.5.\\
&|m(1) - m(0)| \approx 0.231.
\end{aligned}$$

Now consider $\tau_D = 11$, $\tau_R = 9$, so $|\tau_D - \tau_R| = 2 > 1$:

$$\begin{aligned}
&m(11) = \sigma(11) \approx 0.99998, \quad m(9) = \sigma(9) \approx 0.99988.\\
&|m(11) - m(9)| \approx 0.0001.
\end{aligned}$$

So $|\tau_D - \tau_R|$ doubled (from 1 to 2), but $|m(\tau_D) - m(\tau_R)|$ *collapsed* (from 0.231 to 0.0001) because both points moved to the saturated tail of the sigmoid. The variance strictly *decreased* despite a larger trust gap.

Critical Issue

Propositions A1, A2, and A3 all rely on the same false inference: that strict monotonicity of a function implies that the absolute difference of function values is increasing in the absolute difference of arguments. The logistic function saturates at $\pm\infty$, so this fails whenever both arguments are in the same tail.

The main Proposition 1 avoids this error because priors $\pi_D, \pi_R \in (0, 1)$ are bounded, and the mean-preserving spread construction forces one prior up and the other down, which is sufficient for the result on a bounded domain with a strictly monotone function.

Suggestion (fix): The propositions can be repaired in at least two ways:

1. **Add a mean-preserving or symmetric-spread condition:** State that the increase in $|\tau_D - \tau_R|$ occurs via a mean-preserving spread (i.e., $\omega_D\tau_D + \omega_R\tau_R$ is held fixed while $|\tau_D - \tau_R|$ increases). Under this condition, τ_D increases and τ_R decreases (or vice versa), so $m(\tau_D)$ increases and $m(\tau_R)$ decreases, and the gap widens. This directly parallels the proof of Proposition 1.

2. **Bound the domain:** If τ_D, τ_R are restricted to a compact interval $[\underline{\tau}, \bar{\tau}]$ and the spread is defined as τ_D increasing and τ_R decreasing symmetrically, the result holds on the interior.

3. **State the result in terms of the log-odds gap directly:** Proposition A3 could be stated as: "For any $X_D > X_R$, if X_D increases and X_R decreases (or both shift to widen the gap from the same direction), then $|\sigma(X_D) - \sigma(X_R)|$ increases." This is true when one argument moves toward 0 (where $\sigma' = 1/4$, the maximum) and the other moves away.

The simplest and most consistent fix is option (1): impose the same mean-preserving condition used in Proposition 1. This also makes the three appendix propositions structurally parallel to the main result.

3.3 Concavity Claim in Section 3.2

The text states: "$\phi''(\pi) < 0$ for all $\pi \in (0, 1)$ when $p_H > p_L$." Let us verify. We have:

$$\phi(\pi) = \frac{p_H\pi}{p_H\pi + p_L(1-\pi)}.$$

Let $r = p_H/p_L > 1$. Then $\phi(\pi) = r\pi/(r\pi + (1-\pi))$. A calculation gives:

$$\phi'(\pi) = \frac{r}{(r\pi + 1 - \pi)^2}, \qquad \phi''(\pi) = \frac{-2r(r-1)}{(r\pi + 1 - \pi)^3}.$$

Since $r > 1$ and the denominator is positive for $\pi \in (0, 1)$, we have $\phi''(\pi) < 0$. **The concavity claim is correct.**

3.4 Remark on Proposition A3 Nesting

Proposition A3 states that "when $\tau_g = \kappa_g = 1$ for all groups, $\sigma(X_g) = \phi(\pi_g)$." Since $X_g = a(\pi_g) + \tau_g \kappa_g \ell = \log[\pi_g/(1-\pi_g)] + \ell = \log[\pi_g/(1-\pi_g)] + \log(p_H/p_L)$, we have:

$$\sigma(X_g) = \frac{1}{1+\exp(-X_g)} = \frac{1}{1+\frac{1-\pi_g}{\pi_g} \cdot \frac{p_L}{p_H}} = \frac{p_H \pi_g}{p_H \pi_g + p_L(1-\pi_g)} = \phi(\pi_g).$$

The nesting claim is correct.

4 Iteration 4: Empirical Design and Identification

4.1 Is Competitiveness Actually Polarization?

The primary proxy defines "partisan polarization" as $1 - |D_{s,t^*} - R_{s,t^*}|$, which is maximized at 1.0 when the state is perfectly evenly split. The paper uses "polarization" and "competitiveness" interchangeably.

This conflation is problematic. A state with a 50–50 presidential vote share could reflect:

(a) Two large, ideologically entrenched partisan blocs of equal size (true polarization), or

(b) A large population of moderate, persuadable swing voters (low polarization, high competition).

Under scenario (b), the theoretical mechanism (heterogeneous priors driven by partisan identity) is attenuated, because the marginal voters are not ideologically committed. The paper's proxy cannot distinguish between these scenarios.

Major Issue

Suggestion: (a) Explicitly discuss the competitiveness-vs.-polarization distinction. Cite the political science literature on "sorted" vs. "unsorted" electorates [e.g.,][]fiorina2005. (b) Consider a supplementary measure that captures within-state ideological *spread* rather than electoral balance—e.g., county-level vote share *variance* within each state, which distinguishes a state of uniform 50–50 counties (moderate electorate) from a state with deeply red and deeply blue counties (polarized electorate summing to 50–50).

4.2 Weakness of Alternative Proxies

The paper presents three polarization proxies as "triangulating" the latent concept. But the convergence is weak:

- Presidential competitiveness: significant ($p = 0.040$ for $|CAR|$, $p = 0.010$ for $AbVol$).
- Presidential Esteban-Ray: attenuated, insignificant ($p = 0.124$).

- House-district ER: null ($p = 0.893$).
- State DW-NOMINATE: null ($p = 0.522$).
- National DW-NOMINATE: *wrong sign* ($\beta = -0.009$, $p = 0.050$).
- Affective polarization (ANES): null (all $p > 0.40$ in Table 6).

Only the primary presidential competitiveness measure produces significant results. The paper explains each null away individually (low variance for ER, confounding time trends for national DW-NOMINATE, mode shifts for ANES). But the aggregate picture is that *one specific proxy works and five do not*, which raises concerns about specification search.

Critical Issue

An unsympathetic referee will view the robustness battery as showing that the result is fragile rather than robust—it depends on one specific operationalization of "polarization." The paper needs to address this directly rather than dismissing each null individually.

Suggestions:

1. Compute the county-level within-state vote-share variance as a supplementary measure (suggested above). If this measure also produces significant results, it substantially strengthens the case because it captures a different aspect of state-level political heterogeneity.
2. Report the correlation matrix among all proxy measures to show whether they are actually measuring the same latent object. If correlations are low, the "triangulation" framing is misleading.
3. Be more honest in the robustness discussion: rather than explaining away each null, state clearly that only the presidential competitiveness measure produces consistently significant results, and discuss what this implies about the mechanism (is it competitiveness per se, or something correlated with it?).

4.3 Sample Construction and Attrition

The paper starts with 8,621 auditor-change events and arrives at 678—a retention rate of 7.9%. The attrition is:

- 1,720 8-K/A amendments excluded (reasonable).
- 5,641 unmatched to CRSP/Compustat (81.7% of original filings).
- 40 excluded for insufficient trading days.
- Financial firms and utilities excluded (count not stated).

The 82% attrition from CRSP/Compustat matching is extremely high and raises concerns about the representativeness of the final sample. Most auditor changes occur in small, thinly traded, or OTC firms not covered by CRSP/Compustat. The surviving sample is biased toward larger, exchange-listed firms.

Major Issue

Suggestions:

1. Report the number of observations excluded by the financial/utility filter.
2. Discuss whether the results are generalizable only to exchange-listed, CRSP-covered firms, or whether they should be expected to hold more broadly.
3. If possible, report descriptive statistics comparing matched and unmatched filings on observable characteristics (filing year, presence of disagreements, etc.) to assess selection.

4.4 Forward-Filling and Time Variation

The polarization measure is forward-filled from each presidential election to the four subsequent years (2000 → 2001–2004, etc.). This means within each four-year block, the cross-sectional variation in polarization is time-invariant. Time variation comes only from changes across election cycles. With year fixed effects already in the model, the identifying variation for within-state effects (column 6) is extremely limited: it comes from changes in a state's presidential vote margin across at most 5–6 election cycles over 23 years.

Minor Issue

Suggestion: Report how many state × election-cycle cells contain auditor-change events. If many state-cycle cells are empty, the effective identifying variation is even thinner than the 678 observations suggest.

4.5 Measurement Error and Classical Attenuation

The paper argues that measurement error in the headquarters-state proxy attenuates coefficients toward zero, making findings "conservative." This argument assumes classical measurement error (noise uncorrelated with the true variable and with the error term). If the measurement error is systematic—e.g., if swing states are disproportionately states with large urban centers, diverse economies, and different firm compositions—the direction of bias is unknown.

Minor Issue

Suggestion: Qualify the attenuation claim: "Under classical measurement error assumptions, this noise attenuates estimated coefficients. However, if the proxy is systematically correlated with state-level economic characteristics, the direction of bias is indeterminate."

4.6 Endogeneity of Headquarters Location

Firms choose where to locate their headquarters. If firms with particular governance characteristics (e.g., better-governed firms) systematically locate in swing states (which tend to be larger, economically diverse states), the polarization proxy may pick up governance quality rather than political environment. State fixed effects absorb time-invariant differences, but the paper's primary identification comes from cross-state variation.

Minor Issue

Suggestion: Discuss headquarters endogeneity as a limitation. Consider whether the county-level specification (column 9) helps, since county-level variation is more plausibly exogenous to the firm than state-level variation. Alternatively, note that the within-state specification (column 6) identifies only from time-series changes in polarization, which are less susceptible to this concern.

4.7 The Placebo Test Interpretation

The within-firm earnings-announcement placebo ($\beta = 0.000$, $p = 0.924$, $N = 20{,}800$) is well-designed and supports the interpretation channel. However, the paper should note a caveat: earnings announcements differ from auditor-change disclosures on multiple dimensions beyond ambiguity (frequency, magnitude, investor attention, analyst coverage). The null could reflect any of these differences, not specifically the "interpretation vs. precision" distinction.

Minor Issue

Suggestion: Briefly note that the placebo test rules out the *generic amplification* story but cannot definitively attribute the difference to the interpretation channel specifically, because earnings announcements differ from auditor-change disclosures on many dimensions simultaneously.

5 Iteration 5: Redundancy, Notation, and Exposition

5.1 Repeated Passages and Expressions

The manuscript contains substantial redundancy. The following themes are each explained 3–5 times across different sections:

1. **The proxy caveat** ("headquarters-state polarization is a noisy reduced-form proxy for the political-information environment, not a direct measure of investor beliefs"): This appears in the abstract (lines 32–33), the introduction (lines 113–117), Section 5.4 (lines 856–858), the identification subsection (lines 866–874), the retail-investor discussion (lines 876–898), and the conclusion (lines 1205–1208). **Suggestion:** State the caveat fully once in the identification subsection and reference it briefly elsewhere. The current repetition adds approximately 400 words of redundant text.

2. **The channel-agnosticism discussion** ("the same comparative static arises under heterogeneous priors, trust, or precision"): Fully explained in Section 3 (lines 294–316), restated in Section 3.4 (lines 514–534), and repeated in the discussion paragraph below P3 (lines 514–527). **Suggestion:** Consolidate into a single passage in Section 3.

3. **The ambiguity logic**: Developed in Section 3.4 (lines 426–468), restated in H2 (lines 589–604), and repeated in the cross-sectional variables definition (lines 828–837). **Suggestion:** Define the logic once in Section 3.4, state H2 by reference, and define the indicator in Section 5 without repeating the mechanism.

4. **The phrase "the political-information environment"**: Appears 11 times throughout the manuscript. Consider varying the language.

5. **"Politically shaped priors" / "politically shaped interpretation" / "politically shaped inference"**: These near-synonyms appear collectively over 10 times.

Major Issue

Suggestion: A focused editing pass to eliminate redundancy could shorten the paper by 2–3 pages without losing any content. The proxy caveat, channel-agnosticism, and ambiguity logic each need to be stated fully in one canonical location and referenced (not re-explained) elsewhere.

5.2 Notation Consistency

1. **Investor indexing**: The model introduces a "continuum of investors $i \in [0, 1]$" (line 326) but then immediately restricts to two groups $g \in \{D, R\}$. The continuum is never used. **Suggestion:** Drop the continuum notation. State: "A population of investors belongs to one of two groups, $g \in \{D, R\}$, with population shares ω_D and ω_R." This is simpler and honest about the model's structure.

2. π_i **vs.** π_g **vs.** $\pi_{g(i)}$: The text states "$\pi_i \equiv \pi_{g(i)}$; the subscript i indexes individuals and g indexes groups throughout" (lines 333–334). In practice, the paper sometimes writes π_i when it means the group-level prior, creating ambiguity about whether individual-level heterogeneity within groups is allowed. **Suggestion:** Use π_g exclusively in the model, since there is no within-group heterogeneity.

3. ϕ **vs.** σ: The main text defines $\phi(\pi)$ as the Bayesian posterior (Equation 1). The appendix introduces $\sigma(x) = e^x/(1+e^x)$ (line 1327), which is the same function in log-odds space. The paper notes this equivalence (Proposition A3 nesting remark), but the dual notation is unnecessary. **Suggestion:** Either (a) use ϕ throughout and express the appendix results in terms of ϕ, or (b) define σ early in the theory section and use it throughout.

4. **"Polarization" vs. "Competitiveness"**: The primary variable is defined as "partisan competitiveness" (Equation 3), but the paper calls it "polarization" in the regression equation and throughout the results. These are the same variable, but the terminological switching is confusing. **Suggestion:** Choose one term and use it consistently. Since the paper is about polarization, define the variable as "Partisan Polarization" and note its construction as the complement of the margin.

5.3 Structural and Expositional Suggestions

1. **The Hypothesis Development section (Section 4)** is unusually long for a paper that has only two hypotheses and one exploratory prediction. The event-type prediction alone occupies nearly a full page, including a detailed literature review on auditor resignations. If this is auxiliary, it should be shorter; if core, it should be a formal hypothesis. **Suggestion:** Shorten to one paragraph per hypothesis, moving the resignation literature discussion to Section 2.

2. **The identification subsection (Section 5.5)** is approximately 1,400 words—longer than many papers' entire identification discussions. Much of this is hedging and preemptive defense. While the discussion is thoughtful, it is disproportionate to the paper's length. **Suggestion:** Cut to approximately 800 words by consolidating the four confound discussions into a table or a single paragraph with subpoints.

3. **The DW-NOMINATE discussion** (Section 5.3) and the affective polarization discussion (Section 5.4) together span approximately 900 words for proxies that produce null or wrong-signed results. **Suggestion:** Move the construction details to an appendix and keep only a brief description in the main text.

4. **Section labels**: The label "sec:affective" is assigned to Section 2.4 ("Political Polarization and Capital Markets"), which is misleading since only part of that section concerns affective polarization.

5.4 Minor Writing Issues

1. Line 302: "The model thus provides a microfoundation for the broader mechanism described in the introduction, not a claim that priors are the only operative channel." This sentence is somewhat defensive for its position in the theory section. Consider moving to the discussion paragraph below P3.

2. Line 421: "We therefore treat $|CAR[-1,+1]|$ as a motivated complementary outcome." The phrase "motivated complementary outcome" is awkward. Suggested alternative: "a theoretically motivated but secondary outcome."

3. Lines 1004–1006: "The consistency of the coefficient across specifications with and without firm controls suggests that political geography is not merely proxying for firm characteristics." This inference is valid only if the controls are informative about the potential confounds. Consider adding: "...provided the controls capture the relevant firm-level variation."

4. The abstract is 177 words, which is within journal norms but could be tightened. The phrase "a firm may change auditors for routine reasons, rising audit risk, cost or quality considerations, or disagreement over financial reporting" (lines 26–28) duplicates a similar list in the introduction (lines 75–77). **Suggestion:** Shorten the abstract's list to "for various reasons ranging from routine to adversarial."

5. Several footnotes are very long (e.g., the footnote spanning lines 617–717 on presidential vs. congressional returns). Consider incorporating essential content into the text and moving the rest to an appendix.

6 Summary of Actionable Recommendations

The following table prioritizes the recommendations by severity.

Priority	Recommendation
Critical	**Fix the mathematical error in Propositions A1–A3.** Add a mean-preserving or symmetric-spread condition to each proposition. The current statements are false for arbitrary increases in $\lvert\tau_D - \tau_R\rvert$ (or $\lvert\kappa_D - \kappa_R\rvert$, or $\lvert X_D - X_R\rvert$) due to sigmoid saturation.
Critical	**Address the P3 failure for volume honestly.** The ambiguity prediction is contradicted by the $AbVol$ results. Either develop a formal account of why low-ambiguity events generate more disagreement-driven trading, or downgrade P3 from "primary mechanism test" and reframe the second contribution around the event-type split.
Critical	**Confront the weak alternative-proxy results.** Five of six polarization proxies fail. Present the correlation matrix, discuss what this implies for the mechanism, and consider an additional proxy (within-state county vote-share variance).
Major	**Align outcome ordering with theory.** If $AbVol$ is the primary prediction, present it first in all tables, in the abstract, and in the cross-sectional tests. Report both outcomes for every test.
Major	**Distinguish competitiveness from polarization.** Discuss the possibility that 50–50 states reflect moderate rather than polarized electorates, and cite the relevant political science literature.
Major	**Eliminate redundancy.** Consolidate the proxy caveat, channel-agnosticism, and ambiguity logic into single canonical passages. Cut approximately 2–3 pages.
Major	**Provide microfoundation for $\lvert CAR\rvert$.** Add a brief demand-based argument linking posterior dispersion to unsigned price movements, or explicitly state the additional assumptions required.
Minor	**Simplify notation.** Drop the continuum-of-investors notation; use π_g exclusively; choose one symbol (ϕ or σ) for the posterior map; use "polarization" consistently rather than alternating with "competitiveness."
Minor	**Report sample-attrition diagnostics.** Discuss the 82% attrition from CRSP/Compustat matching and its implications for generalizability.
Minor	**Tighten the hypothesis development.** Shorten the event-type prediction or promote it to a hypothesis; shorten the identification discussion to ∼800 words.

Despite these concerns, the paper addresses a genuinely interesting question, the empirical setting is well-chosen, the baseline result is credible, and the placebo test is well-designed. With the revisions suggested above—particularly the proof fix, the honest treatment of the ambiguity results, and the proxy discussion—the paper would be substantially stronger.